

Netzwerktechnik und IT-Netze

Chapter 2: Application Layer

Vorlesung im WS 2016/2017

Bachelor Informatik

(3. Semester)

Prof. Dr. rer. nat. Andreas Fischer


Fakultät für Elektrotechnik, Medientechnik und Informatik

Overview

- Introduction
- Computer Networks and the Internet
- **Application Layer**
 - WWW, Email, DNS, and more
 - Socket programming
- Transport Layer
- Network Layer
- Link Layer
- P2P Networks

Introduction

- A note on the use of these power point slides:
 - All material copyright 1996-2012© J.F Kurose and K.W. Ross, All Rights Reserved
 - Do not copy or distribute this slide set!



*Computer
Networking: A Top
Down Approach*
6th edition
Jim Kurose, Keith Ross
Addison-Wesley
March 2012

Chapter 2: Roadmap

- Principles of network applications
- Web and HTTP
- FTP
- DNS
- Electronic mail
 - SMTP, POP3, IMAP
- Socket programming with UDP and TCP

Chapter 2: application layer

- Our goals:
 - Conceptual, implementation aspects of network application protocols
 - Transport-layer service models
 - Client-server paradigm
 - Peer-to-peer paradigm
 - Learn about protocols by examining popular application-level protocols
 - HTTP
 - FTP
 - SMTP / POP3 / IMAP
 - DNS

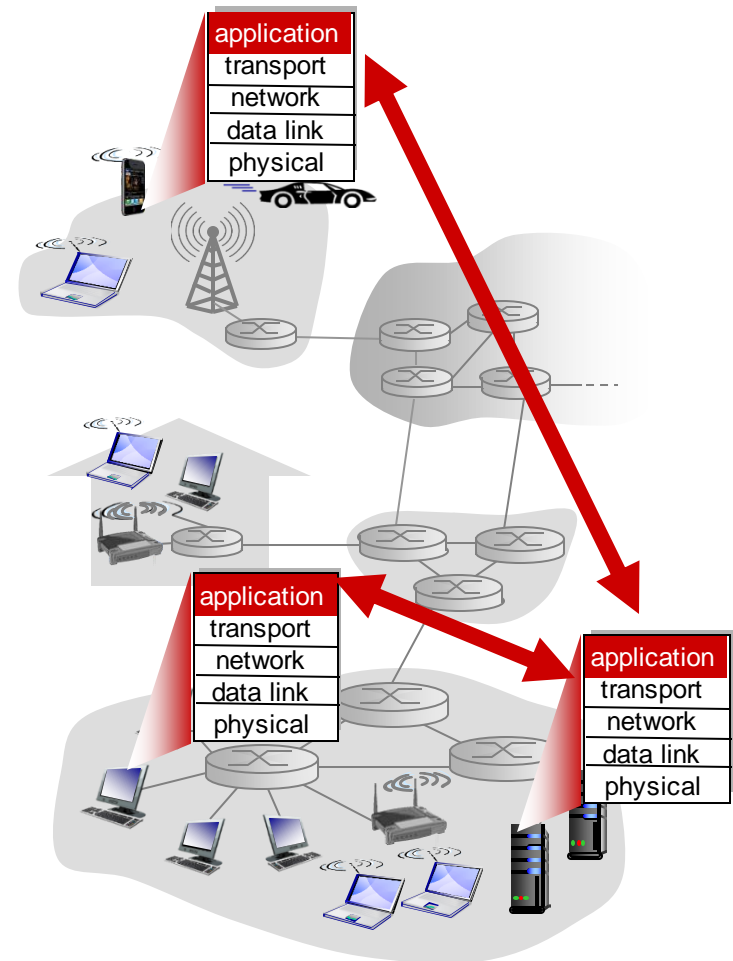
- Our goals:
 - Creating network applications
 - Socket API

Some network apps

- E-mail
- Web
- Text messaging
- Remote login
- P2P file sharing
- Multi-user network games
- Streaming stored video (YouTube, Hulu, Netflix)
- Voice over IP (e.g., Skype)
- Real-time video conferencing
- Social networking
- Search
- ...
- ...

Creating a network app


- Write programs that:
 - Run on (different) end systems
 - Communicate over network
 - e.g., Web server software communicates with browser software
- No need to write software for network-core devices
 - Network-core devices do not run user applications
 - Applications on end systems allows for rapid app development, propagation



Application architectures

- Possible structure of applications:
 - Client-server
 - Peer-to-peer (P2P)


Client-server architecture



- Server:
 - Always-on host
 - Permanent IP address
 - Data centers for scaling
- Clients:
 - Communicate with server
 - May be intermittently connected
 - May have dynamic IP addresses
 - Do not communicate directly with each other

P2P architecture

- No always-on server
- Arbitrary end systems directly communicate
- Peers request service from other peers, provide service in return to other peers
 - Self scalability – new peers bring new service capacity, as well as new service demands
- Peers are intermittently connected and change IP addresses
 - Complex management



Processes communicating

- Process: program running within a host
 - Within same host, two processes communicate using inter-process communication (defined by OS)
 - Processes in different hosts communicate by exchanging messages
 - Aside: applications with P2P architectures have client processes & server processes

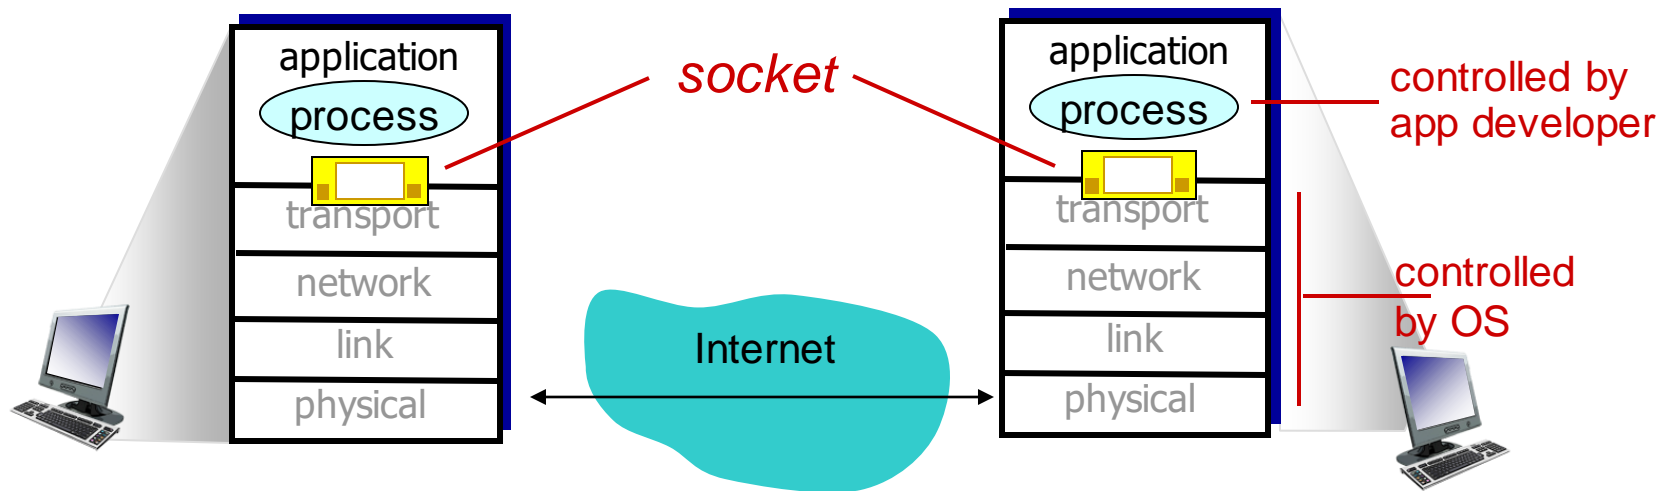
Clients, Servers

Client process: process that initiates communication

Server process: process that waits to be contacted

Sockets

- Process sends/receives messages to/from its **socket**
- Socket analogous to door
 - Sending process shoves message out door
 - Sending process relies on transport infrastructure on other side of door to deliver message to socket at receiving process



Addressing processes

- To receive messages, process must have **identifier**
- Host device has unique 32-bit IP address
- **Q**: does IP address of host on which process runs suffice for identifying the process?
 - **A**: no, many processes can be running on same host
- **Identifier** includes both **IP address** and **port numbers** associated with process on host.
- Example port numbers:
 - **HTTP server: 80**
 - **Mail server: 25**
- To send HTTP message to gaia.cs.umass.edu web server:
 - **IP address: 128.119.245.12**
 - **Port number: 80**
- more shortly...

App-layer protocol defines

- Types of messages exchanged,
 - E.g., request, response
 - Message syntax:
 - What fields in messages & how fields are delineated
 - Message semantics
 - Meaning of information in fields
 - Rules for when and how processes send & respond to messages
- Open protocols:
 - Defined in RFCs
 - Allows for interoperability
 - e.g., HTTP, SMTP
 - Proprietary protocols:
 - E.g., Skype

What transport service does an app need?

- Data integrity
 - Some apps (e.g., file transfer, web transactions) require 100% reliable data transfer
 - Other apps (e.g., audio) can tolerate some loss
- Timing
 - Some apps (e.g., Internet telephony, interactive games) require low delay to be “effective”
- Throughput
 - Some apps (e.g., multimedia) require minimum amount of throughput to be “effective”
 - Other apps (“elastic apps”) make use of whatever throughput they get
- Security
 - Encryption, data integrity, ...

Transport service requirements: common apps

Application	Data loss	Throughput	Time sensitive
file transfer	no loss	elastic	no
e-mail	no loss	elastic	no
Web documents	no loss	elastic	no
real-time audio/video	loss-tolerant	audio: 5kbps-1Mbps video: 10kbps-5Mbps	yes, 100' s msec
stored audio/video	loss-tolerant	same as above	yes, few secs
interactive games	loss-tolerant	few kbps up	yes, 100' s msec
text messaging	no loss	elastic	yes and no

Internet transport protocols services

- TCP service
 - **Reliable transport**: between sending and receiving process
 - **Flow control**: sender won't overwhelm receiver
 - **Congestion control**: throttle sender when network overloaded
 - **Connection-oriented**: setup required between client and server processes
 - **Does not provide**: timing, minimum throughput guarantee, security
- UDP service
 - **Unreliable data transfer**: between sending and receiving process
 - **Does not provide**: reliability, flow control, congestion control, timing, throughput guarantee, security, or connection setup

Internet apps: application, transport protocols

Application	Application layer protocol	Underlying transport protocol
e-mail	SMTP [RFC 2821]	TCP
remote terminal access	Telnet [RFC 854]	TCP
Web	HTTP [RFC 2616]	TCP
file transfer	FTP [RFC 959]	TCP
streaming multimedia	HTTP (e.g., YouTube), RTP [RFC 1889]	TCP or UDP
Internet telephony	SIP, RTP, proprietary (e.g., Skype)	TCP or UDP

Securing TCP

- TCP & UDP
 - No encryption
 - Cleartext passwds sent into socket traverse Internet in cleartext
- SSL
 - Provides encrypted TCP connection
 - Data integrity
 - End-point authentication
- SSL is at app layer
 - Apps use SSL libraries, which “talk” to TCP
- SSL socket API
 - Cleartext passwds sent into socket traverse Internet encrypted
 - See Chapter 7

Chapter 2: Roadmap

- Principles of network applications
- **Web and HTTP**
- FTP
- DNS
- Electronic mail
 - SMTP, POP3, IMAP
- Socket programming with UDP and TCP

Web and HTTP

- First, a review...
 - Web page consists of objects
 - Object can be HTML file, JPEG image, Java applet, audio file,...
 - Web page consists of base HTML-file which includes several referenced objects
 - Each object is addressable by a URL, e.g.,

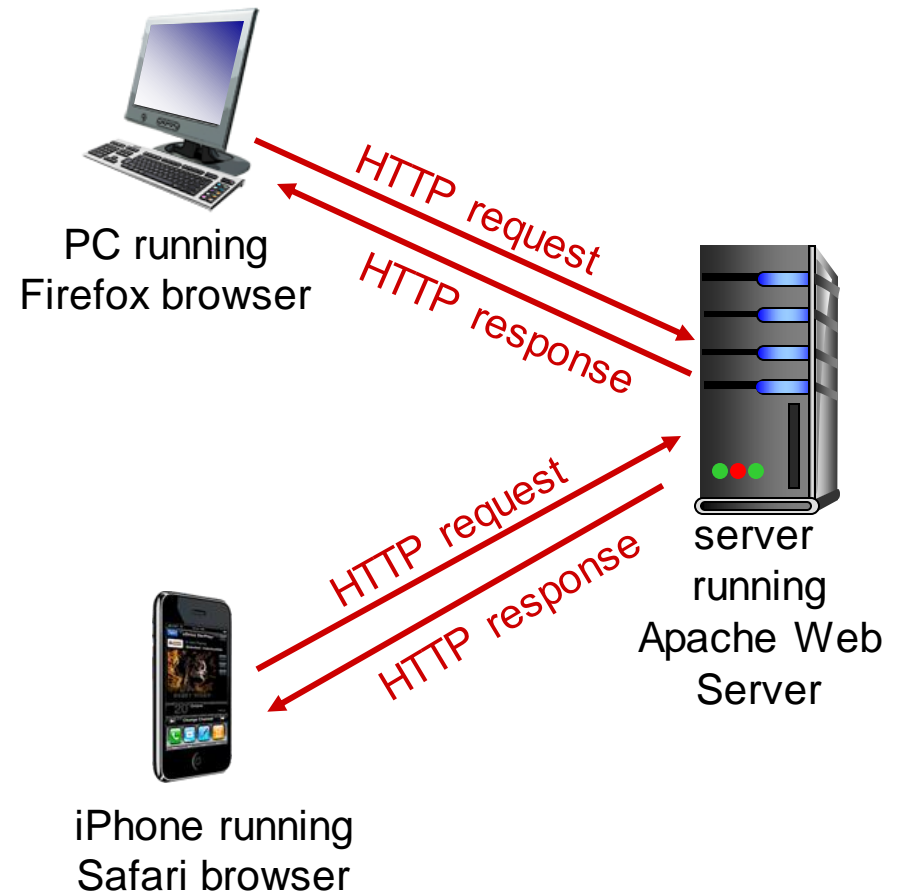
`www.someschool.edu/someDept/pic.gif`

host name

path name

HTTP overview

- HTTP: hypertext transfer protocol
 - Web's application layer protocol
 - Client/server model
 - **Client**: Browser that requests, receives, (using HTTP protocol) and "displays" Web objects
 - **Server**: Web server sends (using HTTP protocol) objects in response to requests



HTTP overview (continued)

- Uses TCP
 - Client initiates TCP connection (creates socket) to server, port 80
 - Server accepts TCP connection from client
 - HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
 - TCP connection closed

- HTTP is “stateless”

- Server maintains no information about past client requests

aside
Protocols that maintain “state” are complex!

- Past history (state) must be maintained
 - If server/client crashes, their views of “state” may be inconsistent, must be reconciled

HTTP connections

- Non-persistent HTTP
 - At most one object sent over TCP connection
 - Connection then closed
 - Downloading multiple objects required multiple connections
- Persistent HTTP
 - Multiple objects can be sent over single TCP connection between client, server

Non-persistent HTTP

- Suppose user enters URL:

- `www.someSchool.edu/someDepartment/home.index`

1a. HTTP client initiates TCP connection to HTTP server (process) at `www.someSchool.edu` on port 80

1b. HTTP server at host `www.someSchool.edu` waiting for TCP connection at port 80. “accepts” connection, notifying client

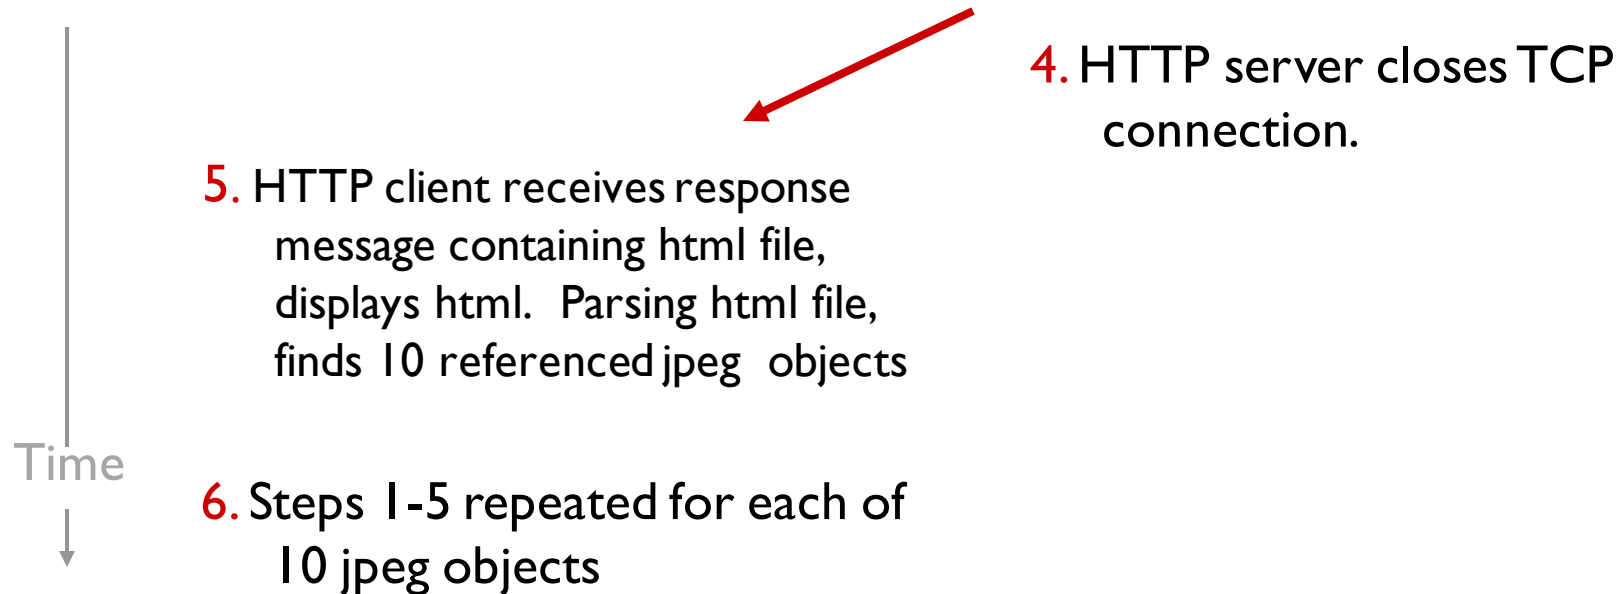
2. HTTP client sends HTTP *request message* (containing URL) into TCP connection socket. Message indicates that client wants object `someDepartment/home.index`

3. HTTP server receives request message, forms *response message* containing requested object, and sends message into its socket

Time




Non-persistent HTTP (cont.)



Non-persistent HTTP: response time

- RTT (round trip time): Time for a small packet to travel from client to server and back
- HTTP response time:
 - One RTT to initiate TCP connection
 - One RTT for HTTP request and first few bytes of HTTP response to return
 - File transmission time
 - Non-persistent HTTP response time = $2\text{RTT} + \text{file transmission time}$



Persistent HTTP

- Non-persistent HTTP issues:
 - Requires 2 RTTs per object
 - OS overhead for each TCP connection
 - Browsers often open parallel TCP connections to fetch referenced objects
- Persistent HTTP:
 - Server leaves connection open after sending response
 - Subsequent HTTP messages between same client/server sent over open connection
 - Client sends requests as soon as it encounters a referenced object

Persistent HTTP

Persistent HTTP **without** pipelining:

- Client sends new request for new objects only, if the response of the previous request has already arrived.
- One RTT for each object.

Persistent **with** pipelining:

- Standard in HTTP/1.1.
- Client sends requests as soon as a reference to an object is found.
- Ideally only a little more than a single RTT for the loading of all referenced objects.

HTTP request message

- Two types of HTTP messages: request, response
- HTTP request message:
 - ASCII (human-readable format)

request line
(GET, POST,
HEAD commands)

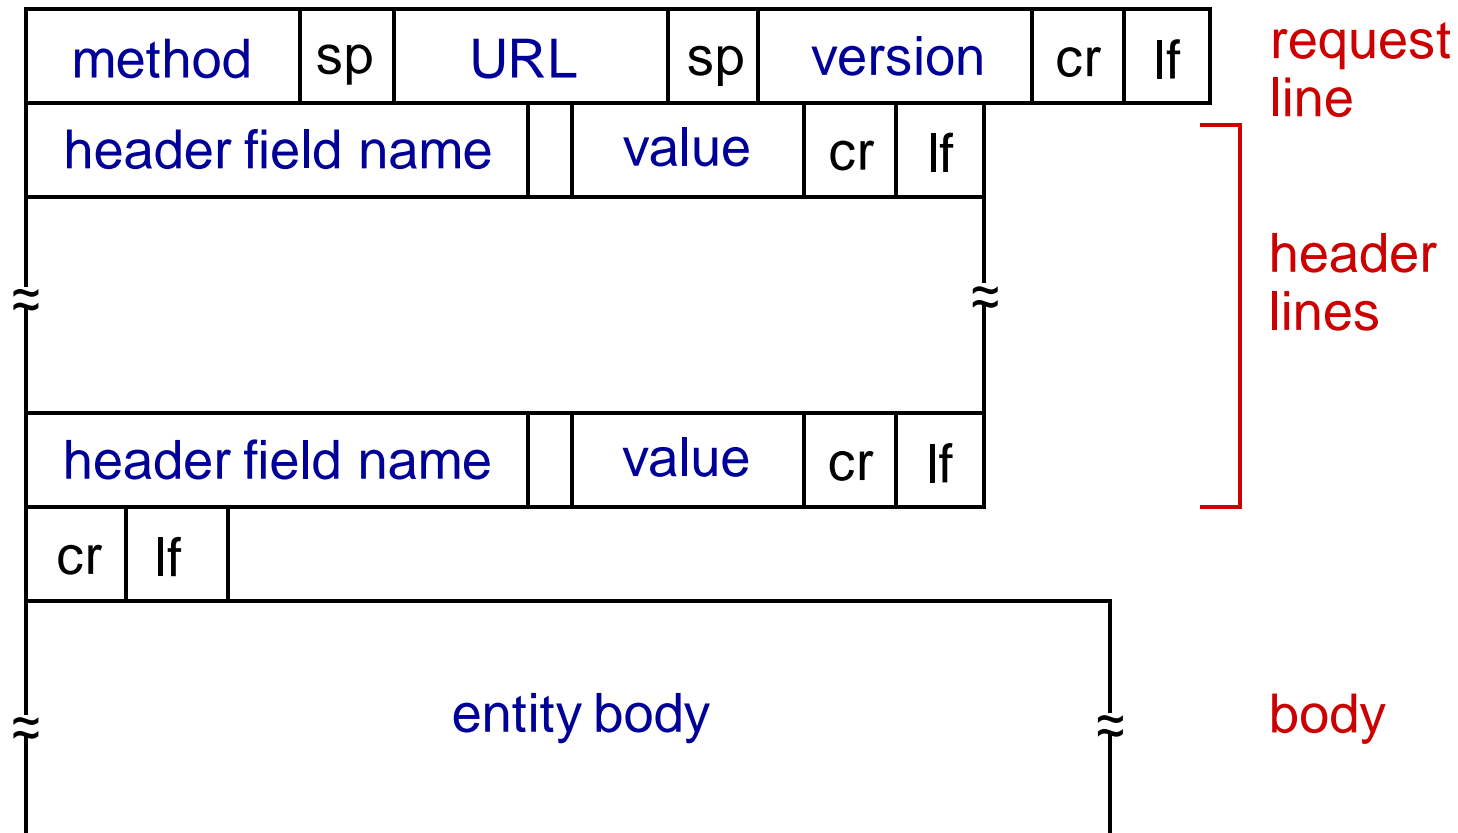
header
lines

carriage return,
line feed at start
of line indicates
end of header lines

```
GET /index.html HTTP/1.1\r\n
Host: www-net.cs.umass.edu\r\n
User-Agent: Firefox/3.6.10\r\n
Accept: text/html,application/xhtml+xml\r\n
Accept-Language: en-us,en;q=0.5\r\n
Accept-Encoding: gzip,deflate\r\n
Accept-Charset: ISO-8859-1,utf-8;q=0.7\r\n
Keep-Alive: 115\r\n
Connection: keep-alive\r\n
\r\n
```

carriage return character
line-feed character

HTTP request message: general format



Uploading form input

- POST method:
 - Web page often includes form input
 - Input is uploaded to server in entity body
- URL method:
 - Uses GET method
 - Input is uploaded in URL field of request line:
 - `www.somesite.com/animalsearch?monkeys&banana`

Method types

- HTTP/1.0:
 - GET
 - POST
 - HEAD
 - Asks server to leave requested object out of response
- HTTP/1.1:
 - GET, POST, HEAD
 - PUT
 - Uploads file in entity body to path specified in URL field
 - DELETE
 - Deletes file specified in the URL field

HTTP response message

status line
(protocol
status code
status phrase)

header
lines

data, e.g.,
requested
HTML file

```
HTTP/1.1 200 OK\r\n
Date: Sun, 26 Sep 2010 20:09:20 GMT\r\n
Server: Apache/2.0.52 (CentOS)\r\n
Last-Modified: Tue, 30 Oct 2007 17:00:02
GMT\r\n
ETag: "17dc6-a5c-bf716880"\r\n
Accept-Ranges: bytes\r\n
Content-Length: 2652\r\n
Keep-Alive: timeout=10, max=100\r\n
Connection: Keep-Alive\r\n
Content-Type: text/html; charset=ISO-8859-1\r\n
\r\n
data data data data data ...
```

HTTP response status codes

- Status code appears in 1st line in server-to-client response message.
 - **200 OK**
 - Request succeeded, requested object later in this msg
 - **301 Moved Permanently**
 - Requested object moved, new location specified later in this msg (Location:)
 - **400 Bad Request**
 - Request msg not understood by server
 - **404 Not Found**
 - Requested document not found on this server
 - **505 HTTP Version Not Supported**

Using HTTP with Telnet

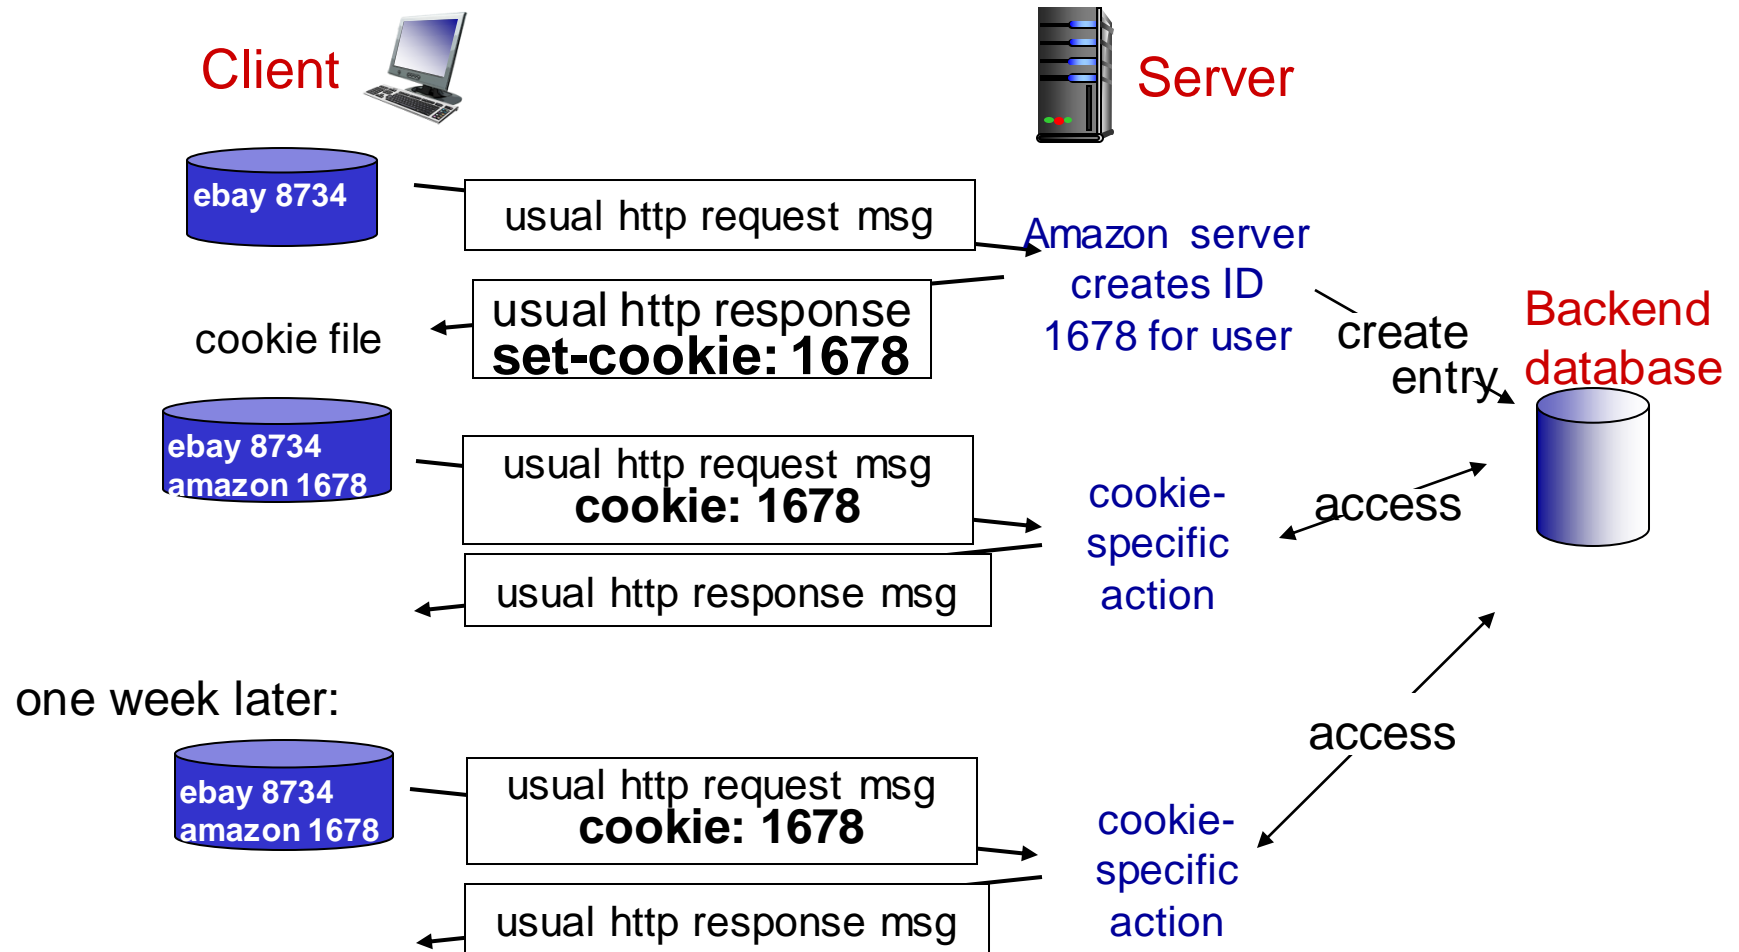
- Download website with “putty” (Telnet-Client):
- Settings: “RAW”, “Never Close”, Port 80, www.bing.com

```
GET / HTTP/1.1\r\n
Host: www.bing.com\r\n
\r\n
```

User-server state: cookies

- Many web sites use cookies
- Four components:
 - Cookie header line of HTTP response message
 - Cookie header line in next HTTP request message
 - Cookie file kept on user's host, managed by user's browser
 - Back-end database at web site
- Example:
 - Susan always accesses Internet from same PC
 - Visits specific e-commerce site for first time
 - When initial HTTP requests arrives at site, site creates:
 - Unique ID
 - Entry in backend database for ID

Cookies: keeping “state” (cont.)



Cookies (continued)

- What cookies can be used for:
 - Authorization
 - Shopping carts
 - Recommendations
 - User session state (web e-mail)
- How to keep “state”:
 - Protocol endpoints: maintain state at sender/receiver over multiple transactions
 - Cookies: HTTP messages carry state


aside

Cookies and privacy:

- Cookies permit sites to learn a lot about you
- You may supply name and e-mail to sites

Web caches (proxy server)

- **Goal:** satisfy client request without involving origin server
 - User sets browser: Web accesses via cache
 - Browser sends all HTTP requests to cache
 - Object in cache: cache returns object
 - Else cache requests object from origin server, then returns object to client




More about Web caching

- Cache acts as both client and server
 - Server for original requesting client
 - Client to origin server
 - Typically cache is installed by ISP (university, company, residential ISP)
- **Why web caching?**
 - Reduce response time for client request
 - Reduce traffic on an institution's access link
 - Internet dense with caches: enables “poor” content providers to effectively deliver content (so too does P2P file sharing)

Caching example


- Assumptions
 - AVG object size: 100K bits
 - AVG request rate from browsers to origin servers: 15/sec
 - AVG data rate to browsers: 1.50 Mbps
 - RTT from institutional router to any origin server: 2 sec
 - access link rate: 1.54 Mbps
- Consequences:
 - LAN utilization: 15%
 - Access link utilization = **99%** *problem!*
 - Total delay = Internet delay + access delay + LAN delay
 = 2 sec + minutes + msecs



Caching example: Fatter access link

- Assumptions
 - AVG object size: 100K bits
 - AVG request rate from browsers to origin servers: 15/sec
 - AVG data rate to browsers: 1.50 Mbps
 - RTT from institutional router to any origin server: 2 sec
 - access link rate: 1.54 Mbps
- Consequences:
 - LAN utilization: 15%
 - Access link utilization = 9,9%
 - Total delay = Internet delay + access delay + LAN delay

15.4 Mbps




Cost: increased access link speed (not cheap!)

Caching example: Install local cache


- Assumptions
 - AVG object size: 100K bits
 - AVG request rate from browsers to origin servers: 15/sec
 - AVG data rate to browsers: 1.50 Mbps
 - RTT from institutional router to any origin server: 2 sec
 - access link rate: 1.54 Mbps
- Consequences:
 - LAN utilization: ?%
 - Access link utilization = ?%
 - Total delay = ?

Cost: web cache (cheap!)




Caching example: Install local cache

- Calculating access link utilization, delay with cache:
 - Suppose cache hit rate is 0.4
 - 40% requests satisfied at cache
 - Access link utilization:
 - 60% of requests use access link
 - Data rate to browsers over access link = $0.6 * 1.50 \text{ Mbps} = 0.9 \text{ Mbps}$
 - utilization = $0.9/1.54 = 58\%$



Caching example: Install local cache

- Calculating access link utilization, delay with cache:
- Total delay
 - $= 0.6 * (\text{delay from origin servers}) + 0.4 * (\text{delay when satisfied at cache})$
 - $= 0.6 (2.01) + 0.4 (\sim \text{msecs})$
 - $= \sim 1.2 \text{ secs}$
 - Less than with 154 Mbps link (and cheaper too!)



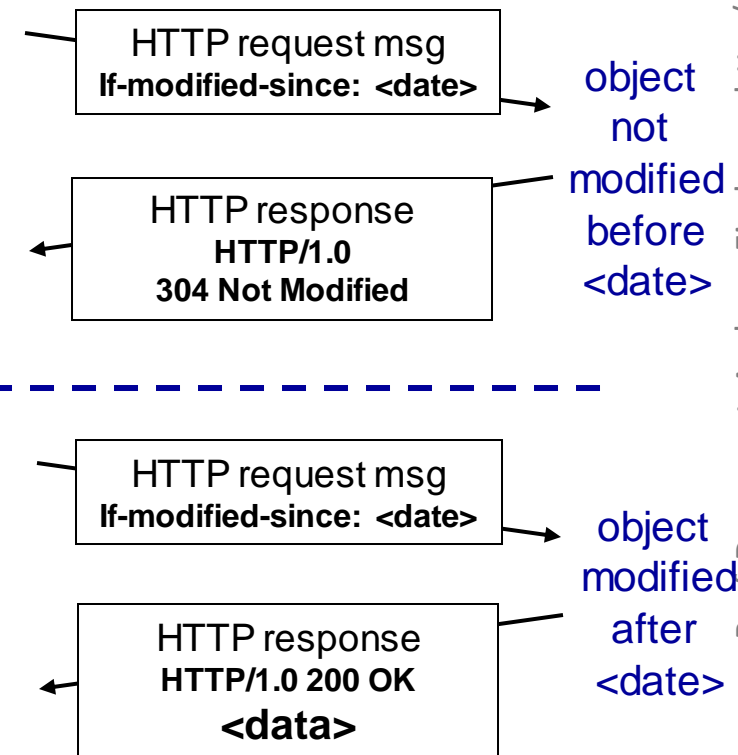
Conditional GET

- Goal: don't send object if cache has up-to-date cached version
 - No object transmission delay
 - Lower link utilization
- Cache: specify date of cached copy in HTTP request
 - If-modified-since: <date>
- Server: response contains no object if cached copy is up-to-date:
 - HTTP/1.0 304 not modified

client



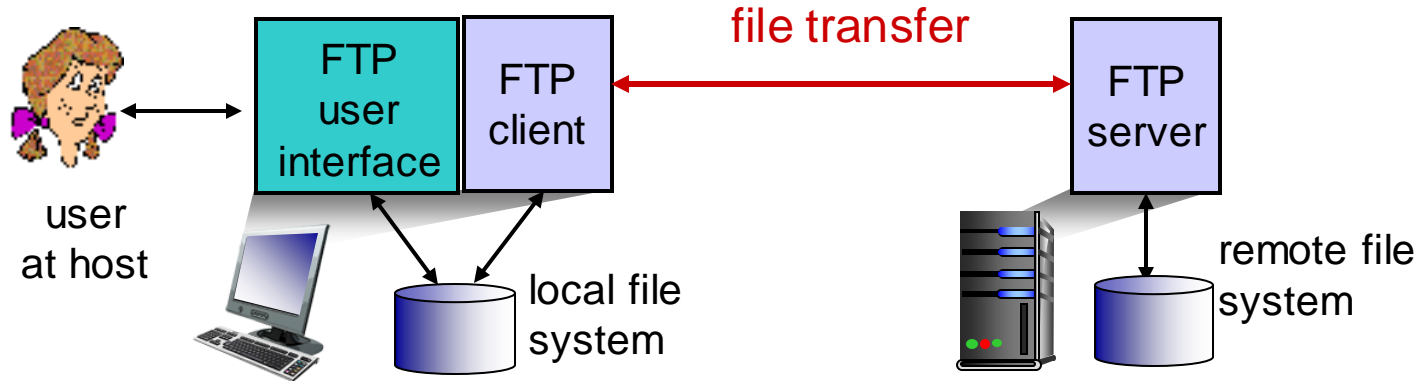
server



Chapter 2: Roadmap

- Principles of network applications
- Web and HTTP
- **FTP**
- DNS
- Electronic mail
 - SMTP, POP3, IMAP
- Socket programming with UDP and TCP


FTP: the file transfer protocol



- Transfer file to/from remote host
- Client/server model
 - client: side that initiates transfer (either to/from remote)
 - server: remote host
- FTP: RFC 959
- FTP server: port 21

FTP: separate control, data connections

- FTP client contacts FTP server at port 21, using TCP
- Client authorized over control connection
- Client browses remote directory, sends commands over control connection
- When server receives file transfer command, **server** opens 2nd TCP data connection (for file) to client
- After transferring one file, server closes data connection



- Server opens another TCP data connection to transfer another file
- Control connection: “**out of band**”
- FTP server maintains “state”: current directory, earlier authentication

FTP commands, responses

- Sample commands
 - Sent as ASCII text over control channel
 - USER username
 - Open FTP-Servers: “anonymous”
 - PASS password
 - Open FTP-Servers: your mail-address
 - LIST return list of file in current directory
 - RETR filename retrieves (gets) file
 - STOR filename stores (puts) file onto remote host
- Sample return codes
 - Status code and phrase (as in HTTP)
 - 331 username OK, password required
 - 125 data connection already open; transfer starting
 - 425 can't open data connection
 - 452 error writing file

FTP commands, responses

- Test FTP-Server: <ftp.avm.de>
- Konsolen ftp:
 - Konsole öffnen: “cmd” (Windows)
 - ftp <ftp.avm.de>
 - Username “anonymous”
 - Password: your mail-address
 - “help” shows a list of all possible instructions
 - dir – lists all files in directory
 - get file – downloads file from server
 - quit – quits ftp
- Open Source Software: FileZilla (Windows/Mac/Linux)

Chapter 2: Roadmap

- Principles of network applications
- Web and HTTP
- FTP
- **DNS**
- Electronic mail
 - SMTP, POP3, IMAP
- Socket programming with UDP and TCP

DNS: domain name system

- People: many identifiers:
 - SSN, name, passport #
 - Internet hosts, routers:
 - IP address (32 bit) - used for addressing datagrams
 - “Name”, e.g.,
www.yahoo.com - used by humans
 - **Q:** How to map between IP address and name, and vice versa ?
- Domain name system:
 - Distributed database implemented in hierarchy of many name servers
 - Application-layer protocol: hosts, name servers communicate to resolve names (address/name translation)
 - Note: core internet function, implemented as application-layer protocol
 - Complexity at network’s “edge”

DNS: services, structure

- DNS services
 - Hostname to IP address translation
 - Host aliasing
 - Canonical, alias names
 - Mail server aliasing
 - Load distribution
 - Replicated web servers: many IP addresses correspond to one name

Why not centralize DNS?

- Single point of failure
- Traffic volume
- Distant centralized database
- Maintenance


A: Doesn't scale!

DNS: a distributed, hierarchical database

Root servers:

TLD servers:

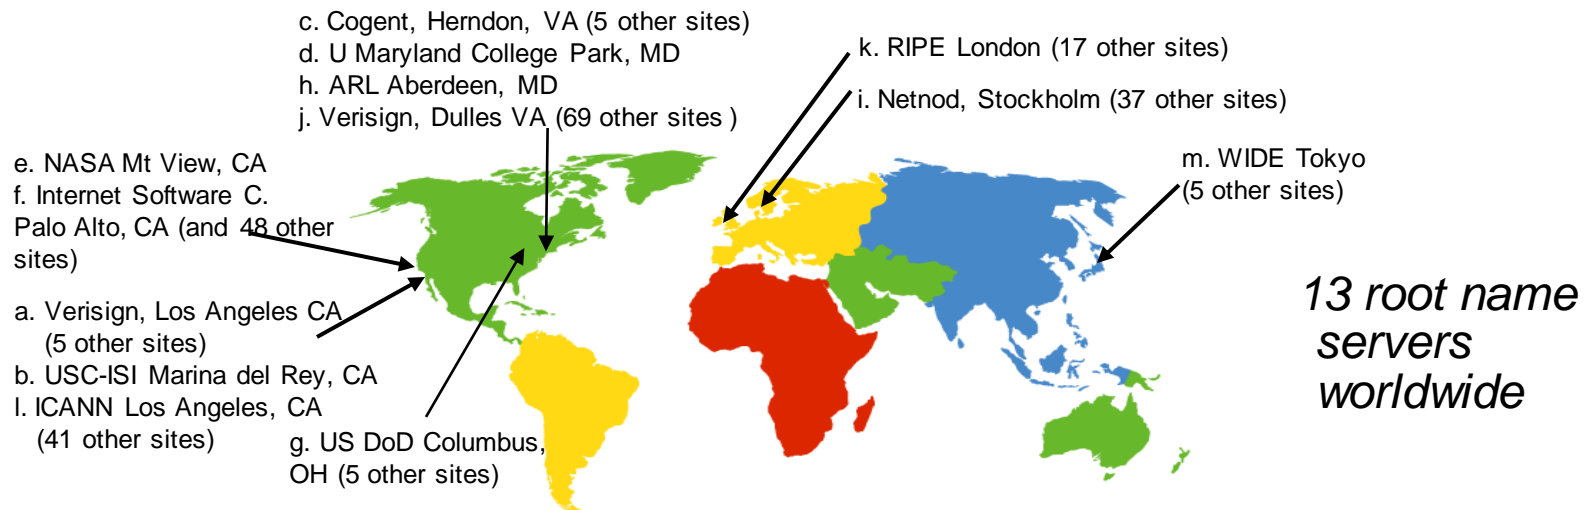
Auth. servers:



- Client wants IP for `www.amazon.com`; 1st approx:
 - Client queries root server to find `.com` DNS server
 - Client queries `.com` DNS server to get `amazon.com` DNS server
 - Client queries `amazon.com` DNS server to get IP address for `www.amazon.com`

DNS: root name servers

- Contacted by local name server that can not resolve name
- Root name server:
 - Contacts authoritative name server if name mapping not known
 - Gets mapping
 - Returns mapping to local name server



TLD, authoritative servers


- Top-level domain (TLD) servers:
 - Responsible for com, org, net, edu, aero, jobs, museums, and all top-level country domains, e.g.: **UK, FR, CA, JP, DE**
 - Network solutions maintains servers for .Com TLD
 - Educause for .Edu TLD
- Authoritative DNS servers:
 - Organization's own DNS server(s), providing authoritative hostname to IP mappings for organization's named hosts
 - Can be maintained by organization or service provider

Local DNS name server

- Does not strictly belong to hierarchy
- Each ISP (residential ISP, company, university) has one
 - Also called “**default name server**”
- When host makes DNS query, query is sent to its local DNS server
 - Has local cache of recent name-to-address translation pairs (but may be out of date!)
 - Acts as proxy, forwards query into hierarchy

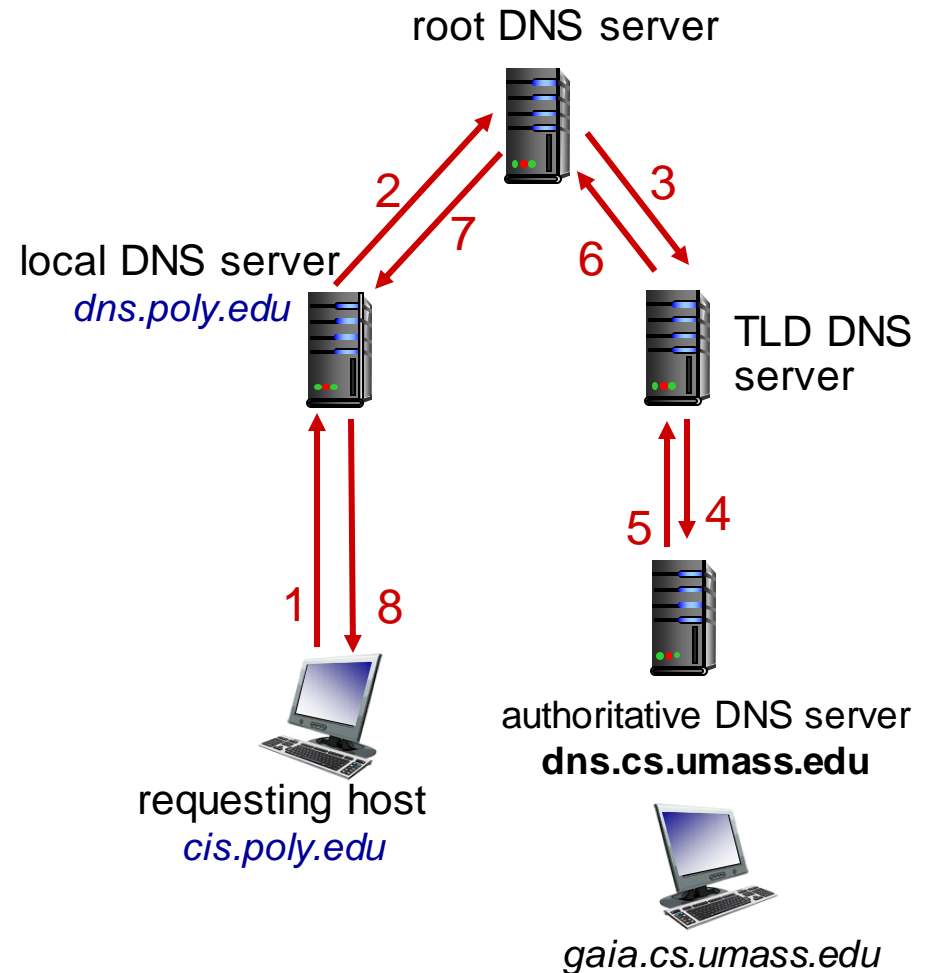
DNS name resolution example

- Host at cis.Poly.Edu wants IP address for gaia.Cs.Umass.Edu
- **Iterated query**
 - Contacted server replies with name of server to contact
 - “I don’t know this name, but ask this server”



DNS name resolution example

- **Recursive query**
 - Puts burden of name resolution on contacted name server
 - Heavy load at upper levels of hierarchy?



DNS: caching, updating records

- Once (any) name server learns mapping, it **caches** mapping
 - Cache entries timeout (disappear) after some time (TTL)
 - TLD servers typically cached in local name servers
 - Thus root name servers not often visited
- Cached entries may be **out-of-date** (best effort name-to-address translation!)
 - If name host changes IP address, may not be known internet-wide until all TTLs expire
- Update/notify mechanisms proposed IETF standard
 - RFC 2136

DNS records

DNS: Distributed DB storing resource records (**RR**)

RR format: (name, value, type, ttl)

type=A

- **name** is hostname
- **value** is IP address

type=NS

- **name** is domain (e.g., foo.com)
- **value** is hostname of authoritative name server for this domain

type=CNAME

- **name** is alias name for some “canonical” (the real) name
- **www.ibm.com** is really **servereast.backup2.ibm.com**
- **value** is canonical name

type=MX

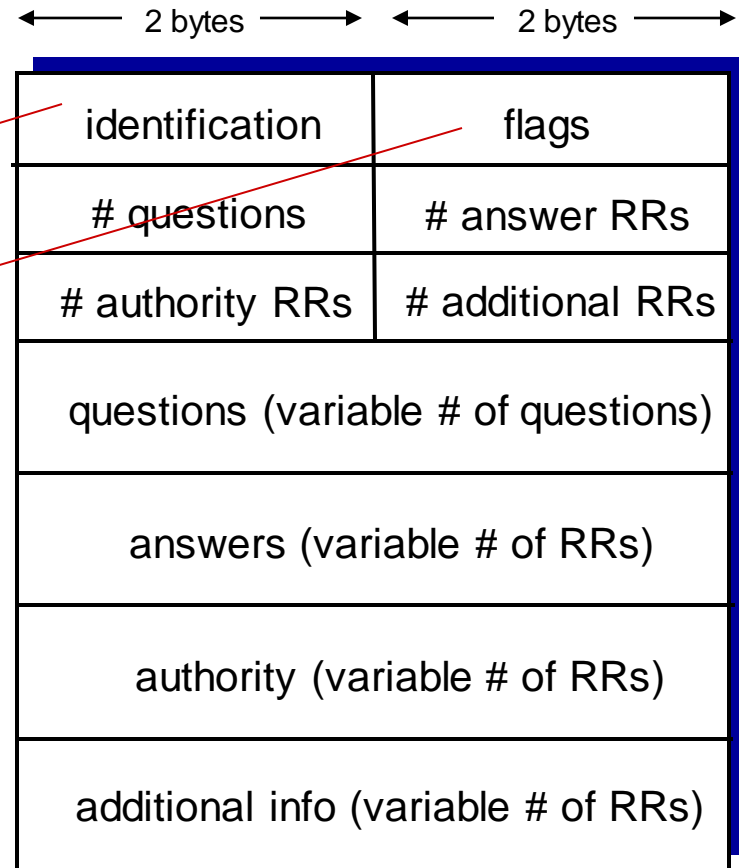
- **value** is name of mailserver associated with **name**

DNS protocol, messages

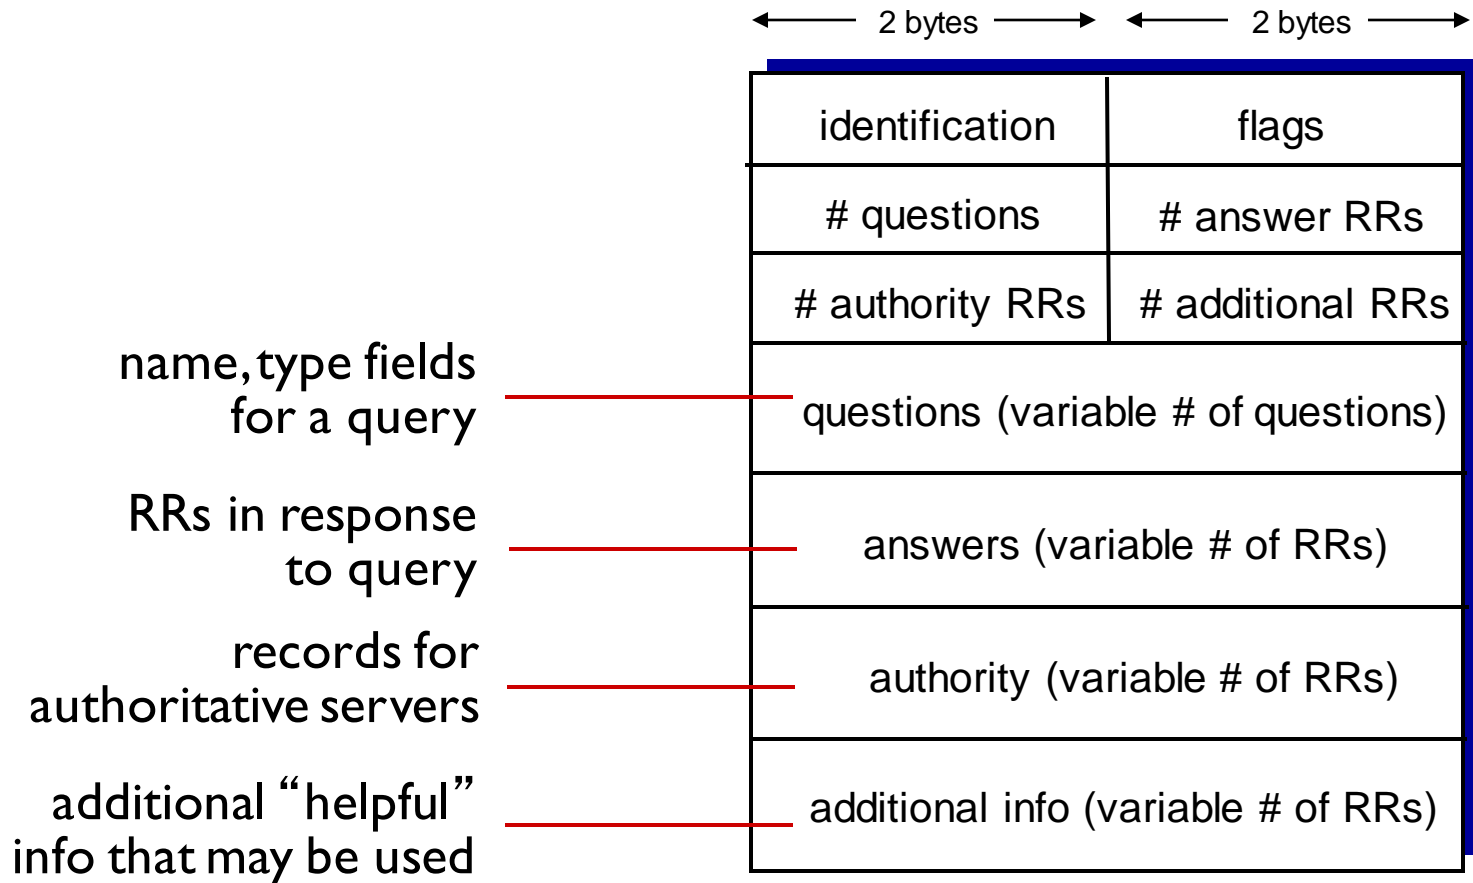
- Query and reply messages, both with same message format

■ Msg header

- Identification: 16 bit # for query, reply to query uses same #
- Flags:
 - Query or reply
 - Recursion desired
 - Recursion available
 - Reply is authoritative



DNS protocol, messages



Inserting records into DNS

- Example: new startup “Network Utopia”
 - Register name networkutopia.com at **DNS registrar** (e.g., Network Solutions)
 - Provide names, IP addresses of authoritative name server (primary and secondary)
 - Registrar inserts two RRs into .com TLD server:
 - (networkutopia.com, dns1.networkutopia.com, NS)
 - (dns1.networkutopia.com, 212.212.212.1, A)

Attacking DNS

- DDoS attacks
 - Bombard root servers with traffic
 - Not successful to date
 - Traffic Filtering
 - Local DNS servers cache IPs of TLD servers, allowing root server bypass
 - Bombard TLD servers
 - Potentially more dangerous
- Redirect attacks
 - Man-in-middle
 - Intercept queries
 - DNS poisoning
 - Send bogus replies to DNS server, which caches
- Exploit DNS for DDoS
 - Send queries with spoofed source address: target IP