

Netzwerktechnik und IT-Netze

Chapter 6: P2P

Vorlesung im WS 2016/2017

Bachelor Informatik

(3. Semester)

Prof. Dr. rer. nat. Andreas Berl


Fakultät für Elektrotechnik, Medientechnik und Informatik

Overview

- Introduction
- Computer Networks and the Internet
- Application Layer
 - WWW, Email, DNS, and more
 - Socket programming
- Transport Layer
- Network Layer
- Link Layer
- **P2P Networks**
- Firewalls

Literatur und Quellen


- A note on the use of these power point slides:
 - All material copyright 1996-2012© J.F Kurose and K.W. Ross, All Rights Reserved
 - Do not copy or distribute this slide set!



*Computer
Networking: A Top
Down Approach*
6th edition
Jim Kurose, Keith Ross
Addison-Wesley
March 2012

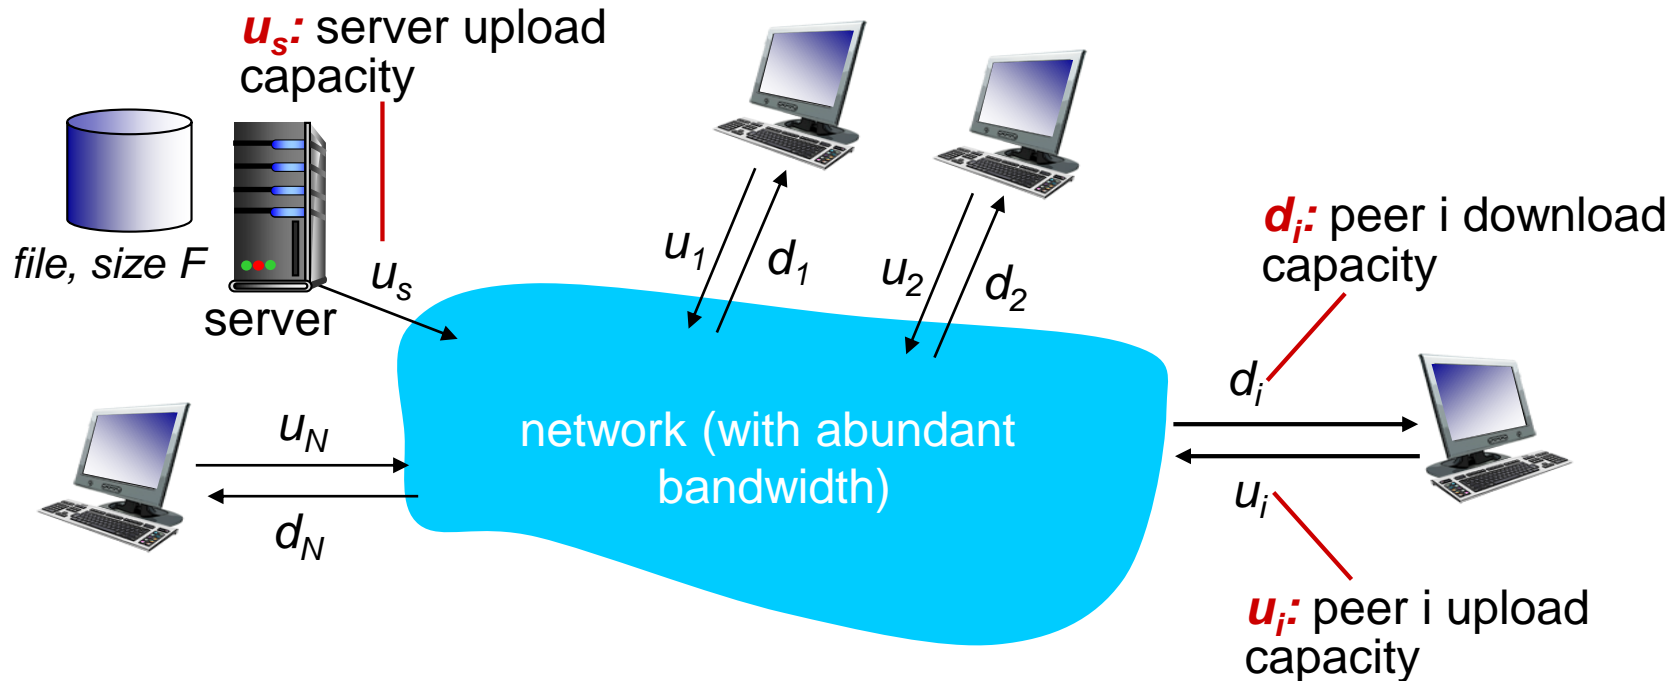
Example: Pure P2P architecture

- No always-on server
- Arbitrary end systems directly communicate
- Peers are intermittently connected and change IP addresses
- Examples:
 - File distribution (BitTorrent)
 - Streaming (Kankan)
 - Voip (Skype)



File distribution: Client-Server vs. P2P


- **Question:** How much time to distribute file (size F) from one server to N peers?
- Peer upload/download capacity is limited resource



File distribution time: Client-server

- **Server transmission:** Must upload N file copies in parallel:
 - Time to send one copy: F/u_s
 - Time to send N copies: NF/u_s
- **Client:** Each client must download file copy
 - d_{\min} = min client download rate
 - min client download time: F/d_{\min}

Time to distribute F to
 N clients using client-
server approach




$$D_{c-s} > \max\{NF/u_s, F/d_{\min}\}$$

Increases linearly in N

File distribution time: P2P

- **Server transmission:** Must upload at least one copy
 - Time to send one copy: F/u_s
- **Client:** Each client must download file copy
 - Min client download time: F/d_{\min}



- **Clients:** As aggregate must download NF bits
 - Max upload rate (limiting max Download rate) is $u_s + \sum u_i$

Time to distribute F to N clients using client-server approach


$$D_{P2P} > \max\{F/u_s, F/d_{\min}, NF/(u_s + \sum u_i)\}$$

Increases linearly in N

... but so does this, as each peer brings service capacity

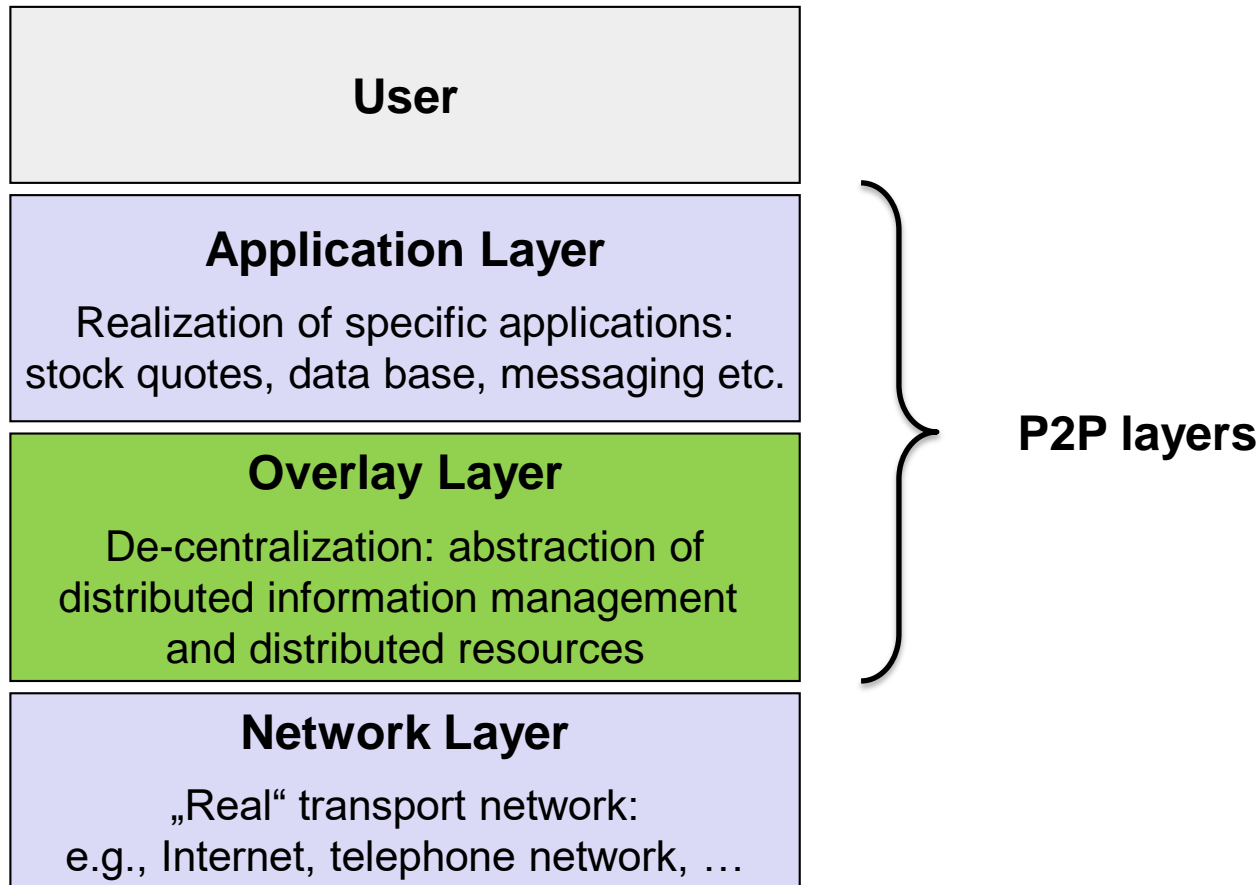
Client-server vs. P2P: Example

- Client upload rate = u , $F/u = 1$ hour, $u_s = 10u$, $d_{\min} \geq u_s$




The Architecture of P2P Systems

- Layering of a P2P system




2.3 A Classification of Peer-to-Peer Systems

- Centralized P2P systems
 - Centralized Server connects Peers
 - Direct interaction between peers
 - Joint usage of distributed resources




- Napster
- ICQ

- Pure P2P systems
 - Completely de-centralized management and usage of resources



- Gnutella
- DHT applications (Past, I3...)


- Hybrid P2P systems
 - Joint usage of distributed resources
 - Direct interaction between peers
 - Some SuperPeers stabilize the Network



- eDonkey
- BitTorrent
- Kazaa

Napster


- First P2P killer application (1999-2001): Napster
 - Exchange of MP3 files



- Central directory server
 - Manages addresses and file lists of peers
 - Does not store files itself

Gnutella

- Decentralized file sharing: Gnutella
 - Host caches provide access points into the P2P network
 - Supports arbitrary file formats
 - Peers are called servents (a combination between server and client)




Gnutella

- Peers are connected via TCP connections
- Search requests are flooded over the Gnutella network
 - TCP broadcast of Ping and Query messages
- Discovery of routing loops through pseudo-unique message IDs (UUID)
 - Temporary storage of UUIDs of received messages
 - Discarding of messages that have been received more than once

Organization of KaZaA

- Organization
 - Each peer is either itself a super node (SN) or is connected to a SN (bandwidth, better reachability)
 - Each SN knows many other SNs (→ mesh overlay) and behaves like a Napster hub (IP addresses, content) from point of view of their „successors“
 - SNs have on average 200-500 successors
 - > 10.000 SNs possible
 - Supernode list server



BitTorrent - Components: Overview

- BitTorrent
- The BT network is partitioned - many isolated swarms, one swarm per content.
- Many decentralized peers:
 - Leechers: Uploading and Downloading
 - Seeders: Only uploading. Typically has a complete copy of the file
- One centralized node per file: Tracker
 - Coordinates the peers of a swarm

BitTorrent - Components: .torrent file

- Meta file used by BT clients to obtain information how to download one file or a batch of files
- Contains:
 - Meta-information about the file (Filename, ...)
 - Resource information (tracker, IP of initial seed, ...)
- One swarm → one .torrent file
- No build-in search. Search is done on web servers which offer .torrent files, for example: <http://thepiratebay.org>

BitTorrent - File deployment

- Process:
 - Generate .torrent file with BT client
 - Register .torrent file at public tracker
 - Alternatively run your own tracker
 - Publish .torrent file on web server or similar
 - Seed the file which should be distributed
- If Tracker dies:
 - Peers can not connect to new peers
 - No new peers can join the swarm
 - Peers still can up- and download to each other
 - High probability that the remaining swarm can finish downloading (when its strongly connected, which is most likely the case)
 - The swarm will split up as peers leave
 - The swarm possibly dies

BitTorrent - File downloading: Overview


- Process overview:
 - User finds the .torrent file referring to the required download
 - Client registers at a tracker and requests other peers of the swarm
 - Client connects to other peers and down- and uploads the content (leecher)
 - After the download is complete, the leecher becomes a seeder, uploading the content



P2P file distribution: BitTorrent


- File divided into 256Kb chunks
- Peers in torrent send/receive file chunks

tracker: tracks peers participating in torrent




Alice arrives ...
... obtains list
of peers from tracker
... and begins exchanging
file chunks with peers in torrent

torrent: group of peers exchanging chunks of a file



P2P file distribution: BitTorrent

- Peer joining torrent:
 - Has no chunks, but will accumulate them over time from other peers
 - Registers with tracker to get list of peers, connects to subset of peers (“neighbors”)




- While downloading, peer uploads chunks to other peers
- Peer may change peers with whom it exchanges chunks
- **Churn**: peers may come and go
- Once peer has entire file, it may (selfishly) leave or (altruistically) remain in torrent

BitTorrent: requesting, sending file chunks

- Requesting chunks:
 - At any given time, different peers have different subsets of file chunks
 - Periodically, Alice asks each peer for list of chunks that they have
 - Alice requests missing chunks from peers, **rarest first**
- Sending chunks: tit-for-tat
 - Alice sends chunks to those four peers currently sending her chunks at highest rate
 - Other peers are choked by Alice (do not receive chunks from her)
 - Re-evaluate top 4 every 10 secs
 - Every 30 secs: randomly select another peer, starts sending chunks
 - “Optimistically unchoke” this peer
 - Newly chosen peer may join top 4

BitTorrent: tit-for-tat

- Alice “optimistically unchokes” Bob
- Alice becomes one of Bob’s top-four providers; Bob reciprocates
- Bob becomes one of Alice’s top-four providers



Principle of Distributed Hash Tables

- Principle of distributed hash tables:
Separation of data & information management
- Nodes have routing informationen (routing tables)
- Routing information provides:
 - Efficient forwarding to „destination“ (content – not location)
 - Statement of existence of content
- Disadvantages:
 - Management of routing information needed
 - Problems with fuzzy requests (e.g, files with name „P*muggle*“)

Principle of Distributed Hash Tables – (cont'd)

- Challenges:
 - (Common P2P requirements) flexibility, reliability and scalability
 - Equal distribution of content to nodes
 - Efficient location of content
 - Permanent adaptation to faults, join, exit of nodes
 - Assignment of responsibilities to new nodes
 - Re-assignment and re-distribution of responsibilities in case of node failure or departure

Operation of Distributed Hash Tables


- Form of operation:
- Mapping of nodes and data into a name space (hash domain)
 - Peers and content are addressed using identifiers (IDs)
 - Joint address space
 - Files are assigned to particular peers
- Storage of / search for data at corresponding nodes
 - Search for data = routing to corresponding node
 - Corresponding node not necessarily known in advance
 - Deterministic statement about availability of data

Mapping into Name Space


- Mapping of nodes:
 - Mapping usually by means of hash functions:
 $H(\text{"IP address"})$ or $H(\text{"host name"})$, respectively
e.g., $H(\text{www.cnn.com}) = 4001$, $H(134.2.11.140) = 4712$
- Mapping of data:
 - Mapping also usually by means of hash functions:
 $H(\text{"file name"})$ or $H(\text{"content"})$, e.g., $H(\text{"movie"}) = 2313$
 - Storage in form of pairs (key, value)
(2313, movie) or (2313, storage location of movie)
 - ID of data can be independent of content (different semantics)

4.1 Operation – Step 1: Mapping into Range

- Step 1:
Mapping of content/nodes into linear range
 - Usually: $0, \dots, 2^m-1 \gg$ number of objects to be stored
 - Mapping of content and nodes into value range by means of a hash function
 - E.g. $H(„/movie/Matrix/divx/en“) \rightarrow 2313$
- Distribution of value range over DHT nodes („responsibilities“)



Value range often represented as a ring

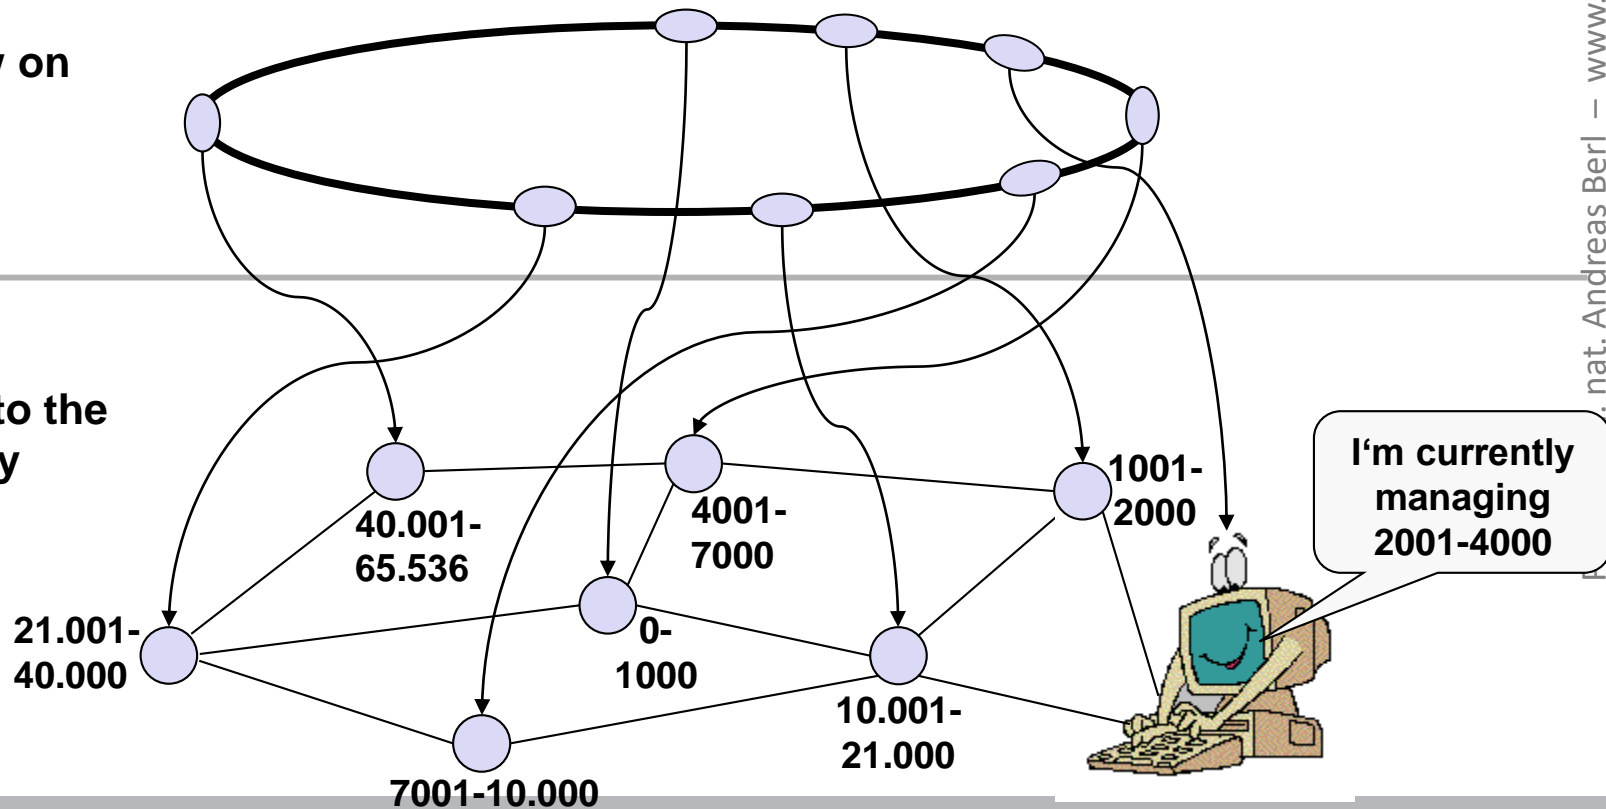


4.1 Distribution of Value Range Across DH Nodes

- Each node is responsible for part of the value range
 - Often highly redundant
 - Repeated adaptation


Logical view on
distributed
hash tables

Mapping onto the
real topology



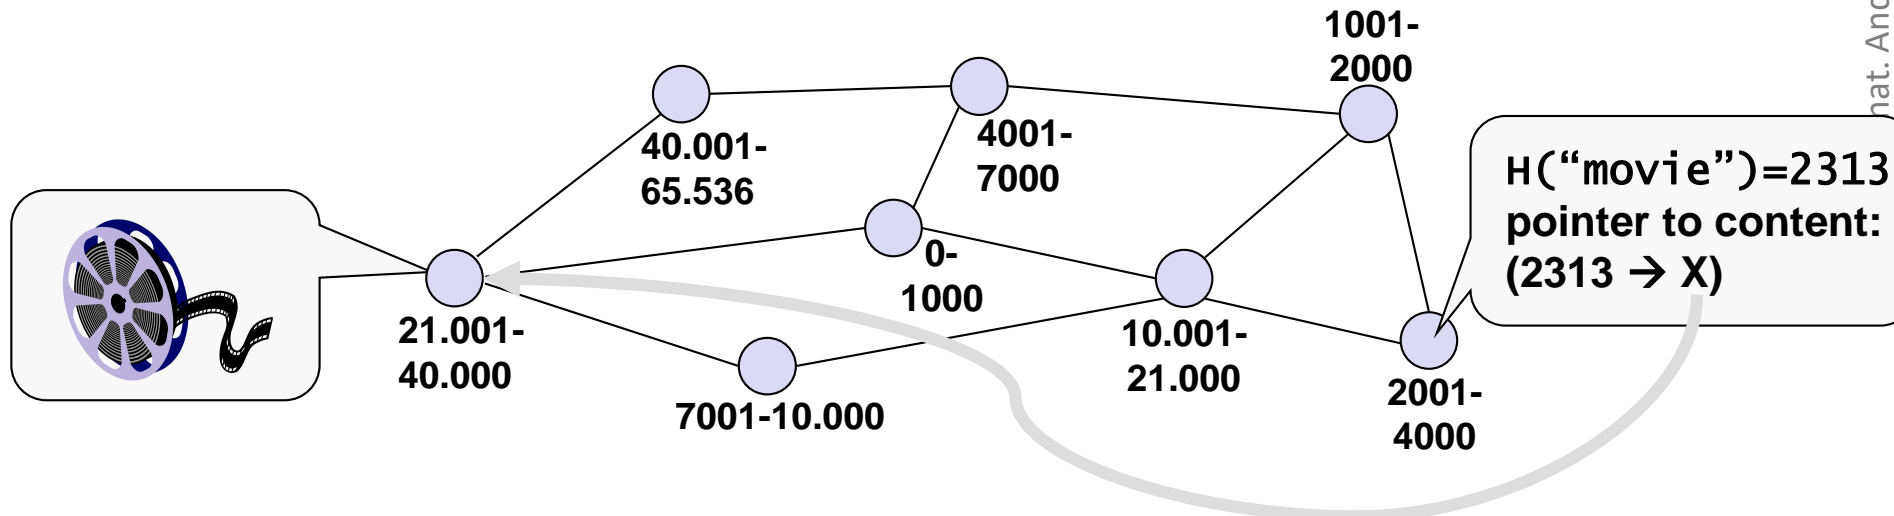
4.1 Storage of Content with DHTs

- How is content stored on the nodes?
 - Assumption: $H(\text{"movie"}) = 2313$ is mapping into value range
- Direct storage:
 - Content is stored on node $H(\text{"movie"})$, i.e, copied to 2313
→ inflexible for large content – o.k. if not too much data



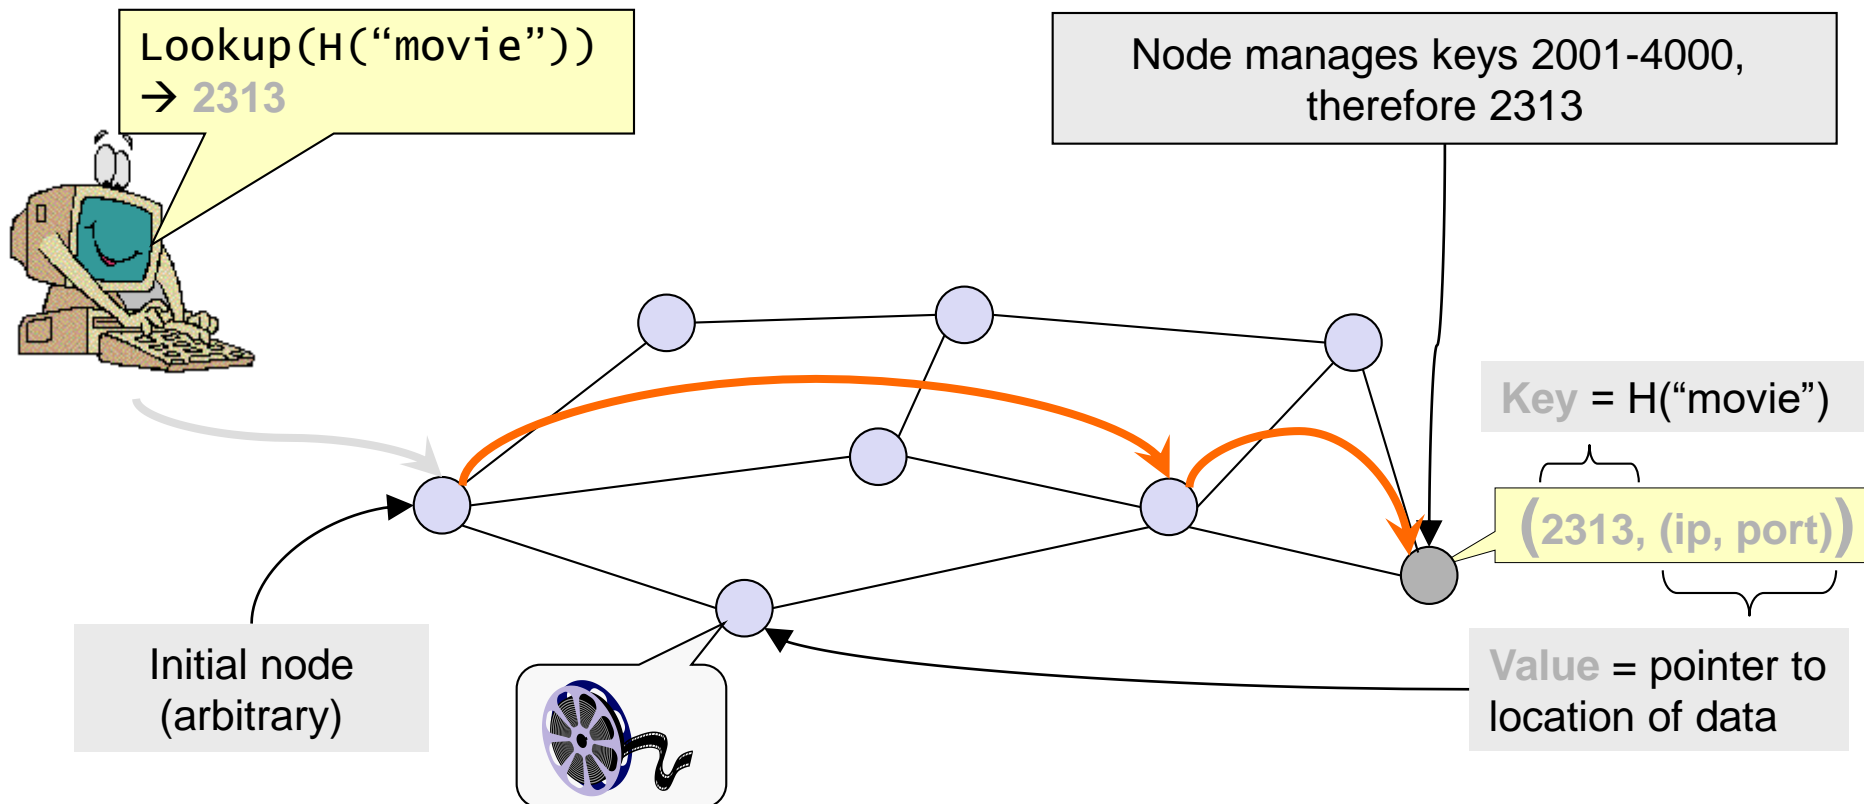
4.1 Storage of Content with DHTs (cont`d)

- Indirect storage:
 - Nodes in DHT store pair (key,value)
 - Key = Hash(„movie”) → 2313
 - Value is often real storage address of content:
(IP, Port) = (134.2.11.140, 4711)
 - More flexible, but one step more to reach content




4.1 Operation – Step 2: Routing to Data (cont'd)

- Search for key
 - Start at arbitrary initial node
 - Routing to requested content (Key)



4.1 Location of Content

- Delivery of information pair (key, value)
- Evaluation and request for real content



4.1 DHT Updating: Joining and Exit of Nodes

- Joining of a new node
 - Assignment of a particular hash range
 - Copying of K/V pairs of hash range (usually with redundancy)
 - Binding into routing environment
- Exit of a node
 - Partitioning of hash range to neighbor nodes
 - Copying of K/V pairs to corresponding nodes
 - Unbinding from routing environment
- Failure of a node
 - Use of redundant K/V pairs (if a node fails)
 - Use of redundant / alternative routing paths
 - Key-value usually still reachable if at least one copy remains

4.1 Resume: Distributed Hash Tables

- Resume: properties of DHTs
 - Use of routing information for efficient search for content
 - Keys are evenly distributed across nodes of DHT (usually with redundancy incorporated)
 - No bottlenecks
 - A continuous increase in number of stored keys is admissible
 - Failure of nodes can be tolerated
 - Survival of attacks possible
 - Self-organizing system
 - Simple and efficient realization
 - Supporting a wide spectrum of applications
 - Flat (hash) key without semantic meaning
 - Value depends on application

DHT Approaches

- Examples of distributed hash tables
 - Chord
UC Berkeley, MIT
 - Pastry
Microsoft Research, Rice University
 - Tapestry
UC Berkeley
 - CAN
UC Berkeley, ICSI
 - P-Grid
EPFL Lausanne
 - Nice, ZigZag
- And many more: Kademlia, Symphony, Viceroy, ...

Chord

- Simple database with (key, value) pairs:
 - Key: movie title; value: IP address
 - Key: human name; value: social security #

Key	Value
John Washington	132-54-3570
Diana Louise Jones	761-55-3791
Xiaoming Liu	385-41-0902
Rakesh Gopal	441-89-1956
Linda Cohen	217-66-5609
.....
Lisa Kobayashi	177-23-0199

Chord: Hash Table

- More convenient to store and search on numerical representation of key
- $\text{Key} = \text{hash}(\text{original key})$

Original Key	Key	Value
John Washington	8962458	132-54-3570
Diana Louise Jones	7800356	761-55-3791
Xiaoming Liu	1567109	385-41-0902
Rakesh Gopal	2360012	441-89-1956
Linda Cohen	5430938	217-66-5609
.....	
Lisa Kobayashi	9290124	177-23-0199

Chord



- Distribute (key, value) pairs over millions of peers
 - Pairs are evenly distributed over peers
- Any peer can query database with a key
 - Database returns value for the key
 - To resolve query, small number of messages exchanged among peers
- Each peer only knows about a small number of other peers
- Robust to peers coming and going (churn)

Chord: Assign key-value pairs to peers

- Rule: assign key-value pair to the peer that has the closest ID.
- Convention: closest is the immediate successor of the key.
- E.g., ID space $\{0, 1, 2, 3, \dots, 63\}$
- Suppose 8 peers: 1, 12, 13, 25, 32, 40, 48, 60
 - If key = 51, then assigned to peer 60
 - If key = 60, then assigned to peer 60
 - If key = 61, then assigned to peer 1


Chord

- Each peer only aware of immediate successor and predecessor.




Chord: Resolving a query

$O(N)$ messages
on average to resolve
query, when there
are N peers




Chord: shortcuts

- Shortcuts
 - Possible to design shortcuts with $O(\log N)$ neighbors, $O(\log N)$ messages in query




Chord: Peer churn



Example: peer 5 abruptly leaves

- Handling peer churn:
 - Peers may come and go (churn)
 - Each peer knows address of its two successors
 - Each peer periodically pings its two successors to check aliveness
 - If immediate successor leaves, choose next successor as new immediate successor

Chord: Peer churn



- Handling peer churn:
 - Peers may come and go (churn)
 - Each peer knows address of its two successors
 - Each peer periodically pings its two successors to check aliveness
 - If immediate successor leaves, choose next successor as new immediate successor

- Peer 4 detects peer 5's departure; makes 8 its immediate successor
- 4 asks 8 who its immediate successor is; makes 8's immediate successor its second successor.

Chord

- Complexity
 - Search overhead: $O(\log N)$
 - Storage requirement per node: $O(\log N)$
 - Fault recovery cost: $O(\log^2 N)$
- Advantages
 - Complexity can be theoretically derived
 - Logarithmic overhead maintained in the face of node failures
- Disadvantages
 - No explicit proximity search for near-by nodes
 - Security has not yet been taken fully into account

Summary

- Centralized, Decentralized and Hybrid P2P
 - Search overhead: $O(\log N)$
 - Storage requirement per node: $O(\log N)$
 - Fault recovery cost: $O(\log^2 N)$
- Advantages
 - Complexity can be theoretically derived
 - Logarithmic overhead maintained in the face of node failures
- Disadvantages
 - No explicit proximity search for near-by nodes
 - Security has not yet been taken fully into account